

Lab Report Knapsack

Simon Jönsson, Fanny Karelius

2017-10-07

Introduction

In this lab we implement different solutions for the **0/1-knapsack**- and **unbounded knapsack** problem. We have implemented both a parallel and non-parallel brute force solution, a dynamic programming solution and a solution using the greedy heuristic. We have documented the runtimes and the profiling of each solution.

Runtime of codes

Bruteforce knapsack solution

Parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:16,], W = 3500, parallel=TRUE))
```

```
##    user  system elapsed  
##  1.042   0.068   1.647
```

Non-parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:16,], W = 3500, parallel=FALSE))
```

```
##    user  system elapsed  
##  0.793   0.038   0.832
```

Dynamic knapsack solution

```
system.time(knapsack_dynamic(x = knapsack_objects[1:500,], W = 3500))
```

```
##    user  system elapsed  
##  5.434   0.081   5.522
```

Greedy knapsack solution

```
system.time(greedy_knapsack(x = knapsack_objects[1:1000000,], W = 3500))
```

```
##    user  system elapsed  
##  1.988   0.156   2.149
```

Profiling

Bruteforce knapsack solution

Parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=TRUE))
```

```
## Reducing depth to 2 (from 8)
```

```
##      time alloc release dups                                ref
## 1 0.003 2.647      0 167      c("matrix", "unlist")
## 2 0.001 0.128      0   0      "matrix"
## 3 0.003 0.053      0  65      c("mclapply", "lapply")
## 4 0.001 0.040      0  13 c("mclapply", "selectChildren")
## 5 0.002 0.056      0  19      c("mclapply", "cleanup")
## 6 0.004 2.293      0   9      c("lapply", "Filter")
##
##      src
## 1 matrix/unlist
## 2 matrix
## 3 mclapply/lapply
## 4 mclapply/selectChildren
## 5 mclapply/cleanup
## 6 lapply/Filter
```

Here we see that apart from allocation, all segments of the code are in similar timesteps - quite tricky to identify bottlenecks. However one could look over using some primitive functions instead of using lapply to find the row with near-optimal value. A suggestions might be using `max()`.

Non-parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=FALSE))
```

```
##      time alloc release dups                                ref      src
## 1 0.006 2.671      0 167 c("matrix", "unlist") matrix/unlist
## 2 0.001 3.572      0   0      "matrix" matrix
## 3 0.001 1.278      0 3813      c("apply", "FUN") apply/FUN
## 4 0.001 1.860      0 1391      "apply" apply
## 5 0.002 4.358      0 3694      c("apply", "FUN") apply/FUN
## 6 0.001 0.000      0 3068      "apply" apply
```

Similar to the parallel solution this code has very little overhead that could be optimized.

Dynamic knapsack solution

```
lineprof(knapsack_dynamic(x = knapsack_objects[1:100,], W = 3500))
```

```
## Reducing depth to 2 (from 70)
```

```
##      time  alloc release  dups                                ref
## 1  0.007   9.915   0.000 5059 c("compiler:::tryCmpfun", "tryCatch")
## 2  0.080 162.500 147.113   65      c("matrix", "replicate")
## 3  0.001   1.228   0.000   0      "matrix"
## 4  0.001   2.656   0.000 2535      character(0)
## 5  0.002   4.049   0.000 7299      "max"
## 6  0.006   8.765   0.000 23657      character(0)
## 7  0.001   1.026   0.000 1003      "max"
## 8  0.001   0.791   0.000 2121      character(0)
## 9  0.002   1.635   0.000 4019      "max"
## 10 0.002   1.431   0.000 2595      character(0)
## 11 0.002   3.048   0.000 6970      "max"
## 12 0.006   8.009   0.000 14581      character(0)
```

```
## 13 0.001    0.638    0.000  2654                "max"
## 14 0.005    3.299    0.000  6937            character(0)
## 15 0.002    1.108    0.000  3169                "max"
## 16 0.003    2.448    0.000  4746            character(0)
## 17 0.001    0.400    0.000   628                "max"
## 18 0.004    1.940    0.000  3446            character(0)
## 19 0.001    0.657    0.000  1387                "max"
## 20 0.007    6.580    0.000 13402            character(0)
## 21 0.001    0.761    0.000  1555                "max"
## 22 0.009   15.841   77.575 36283            character(0)
##                                     src
## 1 compiler:::tryCmpfun/tryCatch
## 2 matrix/replicate
## 3 matrix
## 4
## 5 max
## 6
## 7 max
## 8
## 9 max
## 10
## 11 max
## 12
## 13 max
## 14
## 15 max
## 16
## 17 max
## 18
## 19 max
## 20
## 21 max
## 22
```

Here we identify that the segment in the code that takes most time to run is the replicate function. This could be handled by some other primitive, or maybe the pre-allocation can be circumvented by having dynamic size of the vector.

Greedy knapsack solution

```
lineprof(greedy_knapsack(x = knapsack_objects[1:20000,], W = 3500))
```

```
## Reducing depth to 2 (from 44)
##      time alloc release dups                                ref
## 1 0.007 5.902         0 3083 c("compiler:::tryCmpfun", "tryCatch")
## 2 0.007 3.349         0  173      c("stopifnot", "is.data.frame")
## 3 0.004 7.767         0   44      c("replicate", "sapply")
## 4 0.001 0.395         0    5              "order"
## 5 0.002 0.217         0   31            character(0)
##                                     src
## 1 compiler:::tryCmpfun/tryCatch
## 2 stopifnot/is.data.frame
## 3 replicate/sapply
```

```
## 4 order  
## 5
```

Not alot to improve on here.

Parallelizing brute force knapsack

The performance that could be gained is non-existent since the lapply used doesn't contain any calculations. So there is little to no sequential computations that are done. If we had a computationally heavy lapply segment then we could gain an decrease in computation time.