

Lab Report Knapsack

Simon Jönsson, Fanny Karelius

2017-10-11

Introduction

In this lab we implement different solutions for the **0/1-knapsack**- and **unbounded knapsack** problem. We have implemented both a parallel and non-parallel brute force solution, a dynamic programming solution and a solution using the greedy heuristic. We have documented the runtimes and the profiling of each solution.

Runtime of codes

Bruteforce knapsack solution

Parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:20,], W = 3500, parallel=FALSE))
```

```
##      user  system elapsed  
## 30.299   1.315   31.823
```

Non-parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:20,], W = 3500, parallel=FALSE))
```

```
##      user  system elapsed  
## 19.109   1.474   20.663
```

Dynamic knapsack solution

```
system.time(knapsack_dynamic(x = knapsack_objects[1:500,], W = 3500))
```

```
##      user  system elapsed  
##   5.584   0.234   5.843
```

Greedy knapsack solution

```
system.time(greedy_knapsack(x = knapsack_objects[1:1000000,], W = 3500))
```

```
##      user  system elapsed  
##   1.371   0.196   1.573
```

Profiling

Bruteforce knapsack solution

Parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=TRUE))
```

```
## Reducing depth to 2 (from 16)

##      time alloc release dups                                ref
## 1  0.005 2.647      0  161                                c("matrix", "unlist")
## 2  0.001 0.191      0    0                                "matrix"
## 3  0.004 1.070      0  467                                c("::", "getExportedValue")
## 4  0.001 0.133      0  191                                c("parallel::mclapply", "::")
## 5  0.001 0.087      0    5                                "parallel::mclapply"
## 6  0.003 0.051      0   49                                c("parallel::mclapply", "lapply")
## 7  0.001 0.019      0   11                                "parallel::mclapply"
## 8  0.001 0.035      0    1 c("parallel::mclapply", "lazyLoadDBfetch")
## 9  0.001 0.017      0    8 c("parallel::mclapply", "selectChildren")
##10  0.008 0.485      0   11                                c("parallel::mclapply", "cleanup")
##11  0.004 2.058      0   13                                c("lapply", "Filter")
##12  0.001 0.000      0    9                                c("lapply", "FUN")
##
##                                     src
## 1  matrix/unlist
## 2  matrix
## 3  :/getExportedValue
## 4  parallel::mclapply/::
## 5  parallel::mclapply
## 6  parallel::mclapply/lapply
## 7  parallel::mclapply
## 8  parallel::mclapply/lazyLoadDBfetch
## 9  parallel::mclapply/selectChildren
##10  parallel::mclapply/cleanup
##11  lapply/Filter
##12  lapply/FUN
```

Here we see that apart from allocation, all segments of the code are in similar timesteps - quite tricky to identify bottlenecks. However one could look over using some primitive functions instead of using lapply to find the row with near-optimal value. A suggestions might be using **max()**.

Non-parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=FALSE))
```

```
##      time alloc release dups                                ref                                src
## 1  0.004 2.694      0  161 c("matrix", "unlist") matrix/unlist
## 2  0.001 1.709      0    0                                "matrix" matrix
## 3  0.003 2.292      0 3458 c("apply", "FUN") apply/FUN
## 4  0.001 0.772      0  849                                "apply" apply
## 5  0.009 6.435      0 7837 c("apply", "FUN") apply/FUN
```

Here we see that the apply function and the unlist function creates alot of allocation. Creating a solution which does not use unlist but instead without pre-allocation might help. Also the apply function might be exchanged with using a **max()** primitive. Generating the matrix with given parameters: weight and val one could use **max** to find the maximum val given that the weight $\leq W$.

Dynamic knapsack solution

```
lineprof(knapsack_dynamic(x = knapsack_objects[1:100,], W = 3500))
```

Reducing depth to 2 (from 66)

##	time	alloc	release	dups	ref
## 1	0.011	9.748	0	4879	c("compiler::tryCmpfun", "tryCatch")
## 2	0.206	170.571	0	245	c("matrix", "replicate")
## 3	0.002	0.002	0	0	"matrix"
## 4	0.007	3.931	0	6629	character(0)
## 5	0.002	0.715	0	2395	"max"
## 6	0.002	0.701	0	1025	character(0)
## 7	0.001	0.337	0	997	"max"
## 8	0.001	0.251	0	694	character(0)
## 9	0.001	0.618	0	517	"max"
## 10	0.002	1.035	0	2478	character(0)
## 11	0.001	0.345	0	939	"max"
## 12	0.002	0.869	0	1029	character(0)
## 13	0.001	0.598	0	1478	"max"
## 14	0.002	0.979	0	2367	character(0)
## 15	0.001	0.408	0	891	"max"
## 16	0.001	0.686	0	843	character(0)
## 17	0.002	0.737	0	2814	"max"
## 18	0.002	1.412	0	1608	character(0)
## 19	0.002	1.030	0	2554	"max"
## 20	0.001	0.690	0	1014	character(0)
## 21	0.002	0.869	0	1790	"max"
## 22	0.003	1.410	0	3256	character(0)
## 23	0.005	3.017	0	6000	"max"
## 24	0.001	0.601	0	1322	character(0)
## 25	0.001	0.558	0	1243	"max"
## 26	0.001	0.503	0	1152	character(0)
## 27	0.002	0.432	0	1705	"max"
## 28	0.003	1.173	0	2276	character(0)
## 29	0.001	0.380	0	374	"max"
## 30	0.002	1.083	0	1864	character(0)
## 31	0.001	0.589	0	1158	"max"
## 32	0.001	0.581	0	1218	character(0)
## 33	0.002	0.776	0	1510	"max"
## 34	0.002	0.528	0	1732	character(0)
## 35	0.001	0.167	0	652	"max"
## 36	0.001	0.497	0	343	character(0)
## 37	0.001	0.638	0	1026	"max"
## 38	0.003	1.411	0	2956	character(0)
## 39	0.002	1.072	0	2677	"max"
## 40	0.001	0.714	0	815	character(0)
## 41	0.003	1.559	0	3754	"max"
## 42	0.001	0.687	0	940	character(0)
## 43	0.001	0.628	0	1421	"max"
## 44	0.001	0.554	0	1299	character(0)
## 45	0.002	1.193	0	2545	"max"
## 46	0.001	0.477	0	1064	character(0)
## 47	0.002	0.593	0	1212	"max"
## 48	0.001	0.364	0	996	character(0)
## 49	0.002	1.038	0	2040	"max"
## 50	0.002	1.224	0	1910	character(0)
## 51	0.001	0.336	0	1478	"max"

## 52	0.004	2.463	0 4428	character(0)
## 53	0.002	1.292	0 2783	"max"
## 54	0.001	0.667	0 1241	character(0)
## 55	0.004	2.284	0 4834	"max"
## 56	0.001	0.702	0 1264	character(0)
## 57	0.002	1.208	0 2605	"max"
## 58	0.001	0.671	0 1342	character(0)
## 59	0.002	1.134	0 2649	"max"
## 60	0.001	0.497	0 1083	character(0)
## 61	0.005	2.700	0 5633	"max"
## 62	0.001	0.592	0 974	character(0)
## 63	0.002	1.200	0 2559	"max"
## 64	0.002	1.150	0 2530	character(0)
## 65	0.003	1.479	0 2951	"max"
## 66	0.002	0.807	0 2386	character(0)
## 67	0.005	2.989	0 5599	"max"
## 68	0.002	1.138	0 2381	character(0)
## 69	0.001	0.731	0 925	"max"
## 70	0.001	0.626	0 1509	character(0)
## 71	0.003	1.605	0 3295	"max"
## 72	0.001	0.341	0 1317	character(0)
## 73	0.004	2.171	0 4245	"max"
## 74	0.001	0.443	0 943	character(0)
## 75	0.001	0.000	0 915	"max"
##			src	
## 1	compiler::tryCmpfun/tryCatch			
## 2	matrix/replicate			
## 3	matrix			
## 4				
## 5	max			
## 6				
## 7	max			
## 8				
## 9	max			
## 10				
## 11	max			
## 12				
## 13	max			
## 14				
## 15	max			
## 16				
## 17	max			
## 18				
## 19	max			
## 20				
## 21	max			
## 22				
## 23	max			
## 24				
## 25	max			
## 26				
## 27	max			
## 28				
## 29	max			

```
## 30
## 31 max
## 32
## 33 max
## 34
## 35 max
## 36
## 37 max
## 38
## 39 max
## 40
## 41 max
## 42
## 43 max
## 44
## 45 max
## 46
## 47 max
## 48
## 49 max
## 50
## 51 max
## 52
## 53 max
## 54
## 55 max
## 56
## 57 max
## 58
## 59 max
## 60
## 61 max
## 62
## 63 max
## 64
## 65 max
## 66
## 67 max
## 68
## 69 max
## 70
## 71 max
## 72
## 73 max
## 74
## 75 max
```

Here we identify that the segment in the code that takes most time to run is the replicate function. This could be handled by some other primitive, or maybe the pre-allocation can be circumvented by having dynamic size of the vector.

Greedy knapsack solution

```
lineprof(greedy_knapsack(x = knapsack_objects[1:20000,], W = 3500))
```

```
## Reducing depth to 2 (from 37)
```

```
##      time alloc release dups                                ref
## 1 0.002 2.902         0 3038 c("compiler::tryCmpfun", "tryCatch")
## 2 0.008 1.406         0  218      c("stopifnot", "is.data.frame")
## 3 0.004 9.748         0   44      c("replicate", "sapply")
## 4 0.001 0.336         0    5                                "order"
## 5 0.003 0.622         0   48                                character(0)
##                                     src
## 1 compiler::tryCmpfun/tryCatch
## 2 stopifnot/is.data.frame
## 3 replicate/sapply
## 4 order
## 5
```

Not alot to improve on here.

Parallelizing brute force knapsack

The performance that could be gained is non-existent since the lapply used doesn't contain any calculations. So there is little to no sequential computations that are done. If we had a computationally heavy lapply segment then we could gain an decrease in computation time.