

Lab Report Knapsack

Simon Jönsson, Fanny Karelius

2017-10-08

Introduction

In this lab we implement different solutions for the **0/1-knapsack**- and **unbounded knapsack** problem. We have implemented both a parallel and non-parallel brute force solution, a dynamic programming solution and a solution using the greedy heuristic. We have documented the runtimes and the profiling of each solution.

Runtime of codes

Bruteforce knapsack solution

Parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:16,], W = 3500, parallel=FALSE))
```

```
##      user  system elapsed  
##    1.163    0.046    1.210
```

Non-parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:16,], W = 3500, parallel=FALSE))
```

```
##      user  system elapsed  
##    0.878    0.028    0.909
```

Dynamic knapsack solution

```
system.time(knapsack_dynamic(x = knapsack_objects[1:500,], W = 3500))
```

```
##      user  system elapsed  
##    5.370    0.078    5.455
```

Greedy knapsack solution

```
system.time(greedy_knapsack(x = knapsack_objects[1:1000000,], W = 3500))
```

```
##      user  system elapsed  
##    2.009    0.118    2.131
```

Profiling

Bruteforce knapsack solution

Parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=TRUE))
```

```
## Reducing depth to 2 (from 8)
```

```
##      time alloc release dups                                ref
## 1 0.004 2.647      0 167      c("matrix", "unlist")
## 2 0.001 0.267      0   0      "matrix"
## 3 0.005 0.077      0  78      c("mclapply", "lapply")
## 4 0.001 0.035      0   2 c("mclapply", "lazyLoadDBfetch")
## 5 0.001 0.017      0   8 c("mclapply", "selectChildren")
## 6 0.005 0.423      0   6      c("mclapply", "cleanup")
## 7 0.003 1.866      0  17      c("lapply", "Filter")
##                                     src
## 1 matrix/unlist
## 2 matrix
## 3 mclapply/lapply
## 4 mclapply/lazyLoadDBfetch
## 5 mclapply/selectChildren
## 6 mclapply/cleanup
## 7 lapply/Filter
```

Here we see that apart from allocation, all segments of the code are in similar timesteps - quite tricky to identify bottlenecks. However one could look over using some primitive functions instead of using lapply to find the row with near-optimal value. A suggestions might be using `max()`.

Non-parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=FALSE))
```

```
## Reducing depth to 2 (from 3)
```

```
##      time alloc release dups                                ref                                src
## 1 0.006 2.684      0 167 c("matrix", "unlist") matrix/unlist
## 2 0.001 1.559      0   0      "matrix" matrix
## 3 0.004 3.208      0 4254      c("apply", "FUN") apply/FUN
## 4 0.001 0.784      0  860      "apply" apply
## 5 0.003 2.503      0 2704      c("apply", "FUN") apply/FUN
## 6 0.001 0.798      0  872      "apply" apply
## 7 0.002 1.564      0 1696      c("apply", "FUN") apply/FUN
## 8 0.001 0.786      0  872      "apply" apply
## 9 0.001 0.000      0  854      c("apply", "FUN") apply/FUN
```

Similar to the parallel solution this code has very little overhead that could be optimized.

Dynamic knapsack solution

```
lineprof(knapsack_dynamic(x = knapsack_objects[1:100,], W = 3500))
```

```
## Reducing depth to 2 (from 66)
```

```
##      time   alloc release  dups                                ref
## 1  0.011   9.914   0.000  4997 c("compiler::tryCmpfun", "tryCatch")
## 2  0.088 162.834 147.201   127      c("matrix", "replicate")
## 3  0.001   0.676   0.000    0      "matrix"
## 4  0.001   6.027   0.000  1396      "max"
## 5  0.002   6.064   0.000 23529      character(0)
```

## 6	0.001	2.700	0.000	1451	"max"
## 7	0.005	10.313	0.000	23615	character(0)
## 8	0.001	0.786	0.000	3274	"max"
## 9	0.001	0.283	0.000	1626	character(0)
## 10	0.003	1.790	0.000	2923	"max"
## 11	0.006	4.878	0.000	10681	character(0)
## 12	0.001	0.760	0.000	756	"max"
## 13	0.002	1.535	0.000	3277	character(0)
## 14	0.002	1.340	0.000	2796	"max"
## 15	0.005	2.938	0.000	6040	character(0)
## 16	0.001	0.780	0.000	1465	"max"
## 17	0.001	0.180	0.000	1613	character(0)
## 18	0.001	0.747	0.000	369	"max"
## 19	0.001	0.812	0.000	1544	character(0)
## 20	0.002	1.579	0.000	3398	"max"
## 21	0.001	0.806	0.000	1544	character(0)
## 22	0.002	1.547	0.000	3158	"max"
## 23	0.001	0.632	0.000	1706	character(0)
## 24	0.001	0.612	0.000	1305	"max"
## 25	0.001	0.825	0.000	1264	character(0)
## 26	0.002	1.334	0.000	2758	"max"
## 27	0.004	2.913	0.000	6046	character(0)
## 28	0.004	5.338	72.305	15335	"max"
## 29	0.002	2.969	0.000	16118	character(0)
##			src		
## 1	compiler::tryCmpfun/tryCatch				
## 2	matrix/replicate				
## 3	matrix				
## 4	max				
## 5					
## 6	max				
## 7					
## 8	max				
## 9					
## 10	max				
## 11					
## 12	max				
## 13					
## 14	max				
## 15					
## 16	max				
## 17					
## 18	max				
## 19					
## 20	max				
## 21					
## 22	max				
## 23					
## 24	max				
## 25					
## 26	max				
## 27					
## 28	max				
## 29					

Here we identify that the segment in the code that takes most time to run is the replicate function. This could be handled by some other primitive, or maybe the pre-allocation can be circumvented by having dynamic size of the vector.

Greedy knapsack solution

```
lineprof(greedy_knapsack(x = knapsack_objects[1:20000,], W = 3500))
```

```
## Reducing depth to 2 (from 46)
```

```
##      time alloc release dups                                ref
## 1 0.003 4.842         0 2736 c("compiler::tryCmpfun", "tryCatch")
## 2 0.006 1.339         0  520      c("stopifnot", "is.data.frame")
## 3 0.003 9.777         0   43      c("replicate", "sapply")
## 4 0.001 0.030         0    6              "order"
## 5 0.004 0.879         0   46      character(0)
##
##      src
## 1 compiler::tryCmpfun/tryCatch
## 2 stopifnot/is.data.frame
## 3 replicate/sapply
## 4 order
## 5
```

Not alot to improve on here.

Parallelizing brute force knapsack

The performance that could be gained is non-existent since the lapply used doesn't contain any calculations. So there is little to no sequential computations that are done. If we had a computationally heavy lapply segment then we could gain an decrease in computation time.