

Lab Report Knapsack

Simon Jönsson, Fanny Karelius

2017-10-09

Introduction

In this lab we implement different solutions for the **0/1-knapsack**- and **unbounded knapsack** problem. We have implemented both a parallel and non-parallel brute force solution, a dynamic programming solution and a solution using the greedy heuristic. We have documented the runtimes and the profiling of each solution.

Runtime of codes

Bruteforce knapsack solution

Parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:16,], W = 3500, parallel=FALSE))
```

```
##      user  system elapsed  
##    1.193    0.053    1.249
```

Non-parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:16,], W = 3500, parallel=FALSE))
```

```
##      user  system elapsed  
##    0.908    0.040    0.952
```

Dynamic knapsack solution

```
system.time(knapsack_dynamic(x = knapsack_objects[1:500,], W = 3500))
```

```
##      user  system elapsed  
##    5.482    0.098    5.609
```

Greedy knapsack solution

```
system.time(greedy_knapsack(x = knapsack_objects[1:1000000,], W = 3500))
```

```
##      user  system elapsed  
##    2.156    0.175    2.369
```

Profiling

Bruteforce knapsack solution

Parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=TRUE))
```

```
## Reducing depth to 2 (from 9)

##      time alloc release dups      ref
## 1  0.008 2.671      0 167      c("matrix", "unlist")
## 2  0.001 0.129      0   0      "matrix"
## 3  0.001 0.138      0  17      c("mclapply", "detectCores")
## 4  0.004 0.062      0  52      c("mclapply", "lapply")
## 5  0.001 0.027      0  11 c("mclapply", "lazyLoadDBfetch")
## 6  0.001 0.003      0   2      c("mclapply", "unserialize")
## 7  0.001 0.038      0   1      "mclapply"
## 8  0.010 0.829      0  13      c("mclapply", "cleanup")
## 9  0.003 1.732      0  13      c("lapply", "Filter")
## 10 0.001 0.000      0  44      character(0)
##
##      src
## 1  matrix/unlist
## 2  matrix
## 3  mclapply/detectCores
## 4  mclapply/lapply
## 5  mclapply/lazyLoadDBfetch
## 6  mclapply/unserialize
## 7  mclapply
## 8  mclapply/cleanup
## 9  lapply/Filter
## 10
```

Here we see that apart from allocation, all segments of the code are in similar timesteps - quite tricky to identify bottlenecks. However one could look over using some primitive functions instead of using lapply to find the row with near-optimal value. A suggestions might be using `max()`.

Non-parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=FALSE))
```

```
## Reducing depth to 2 (from 3)

##      time alloc release dups      ref      src
## 1  0.008 2.828      0 167 c("matrix", "unlist") matrix/unlist
## 2  0.001 1.678      0  16 c("apply", "aperm") apply/aperm
## 3  0.002 1.400      0 2575      "apply" apply
## 4  0.011 7.292      0 7967      c("apply", "FUN") apply/FUN
## 5  0.001 0.623      0  712      "apply" apply
## 6  0.001 0.000      0  676      c("apply", "FUN") apply/FUN
```

Similar to the parallel solution this code has very little overhead that could be optimized.

Dynamic knapsack solution

```
lineprof(knapsack_dynamic(x = knapsack_objects[1:100,], W = 3500))
```

```
## Reducing depth to 2 (from 66)

##      time  alloc release dups      ref
## 1  0.009   9.747   0.000 4489 c("compiler:::tryCmpfun", "tryCatch")
## 2  0.130 168.612 152.813  635      c("matrix", "replicate")
```

## 3	0.002	0.001	0.000	0	"matrix"
## 4	0.001	0.004	0.000	0	character(0)
## 5	0.001	0.605	0.000	7	"max"
## 6	0.009	5.848	0.000	11189	character(0)
## 7	0.003	1.474	0.000	3912	"max"
## 8	0.003	2.055	0.000	4309	character(0)
## 9	0.001	0.796	0.000	1216	"max"
## 10	0.005	2.534	0.000	5231	character(0)
## 11	0.003	2.002	0.000	4920	"max"
## 12	0.007	2.664	0.000	5617	character(0)
## 13	0.001	0.599	0.000	751	"max"
## 14	0.003	1.933	0.000	3515	character(0)
## 15	0.003	1.192	0.000	3358	"max"
## 16	0.001	0.574	0.000	824	character(0)
## 17	0.001	0.457	0.000	1187	"max"
## 18	0.005	2.778	0.000	5365	character(0)
## 19	0.001	0.709	0.000	1318	"max"
## 20	0.010	3.324	0.000	7573	character(0)
## 21	0.003	1.923	0.000	3249	"max"
## 22	0.001	0.319	0.000	1485	character(0)
## 23	0.002	1.117	0.000	1857	"max"
## 24	0.003	1.067	0.000	2355	character(0)
## 25	0.001	0.670	0.000	958	"max"
## 26	0.005	2.500	0.000	5528	character(0)
## 27	0.001	0.386	0.000	1025	"max"
## 28	0.003	1.688	0.000	2656	character(0)
## 29	0.002	1.218	0.000	2958	"max"
## 30	0.003	1.572	0.000	3059	character(0)
## 31	0.001	0.676	0.000	1374	"max"
## 32	0.001	0.761	0.000	1397	character(0)
## 33	0.001	0.722	0.000	1574	"max"
## 34	0.001	0.483	0.000	1493	character(0)
## 35	0.002	1.366	0.000	2518	"max"
## 36	0.010	6.026	0.000	12892	character(0)
## 37	0.001	0.554	0.000	865	"max"
## 38	0.002	0.768	0.000	2329	character(0)
## 39	0.002	0.000	75.874	404	"max"
## 40	0.001	1.043	0.000	5231	character(0)
## 41	0.001	0.305	0.000	2155	"max"
## 42	0.002	0.600	0.000	1417	character(0)
## 43	0.001	1.174	0.000	451	"max"
## 44	0.015	10.818	0.000	24781	character(0)
##				src	
## 1	compiler:::tryCmpfun/tryCatch				
## 2	matrix/replicate				
## 3	matrix				
## 4					
## 5	max				
## 6					
## 7	max				
## 8					
## 9	max				
## 10					
## 11	max				

```

## 12
## 13 max
## 14
## 15 max
## 16
## 17 max
## 18
## 19 max
## 20
## 21 max
## 22
## 23 max
## 24
## 25 max
## 26
## 27 max
## 28
## 29 max
## 30
## 31 max
## 32
## 33 max
## 34
## 35 max
## 36
## 37 max
## 38
## 39 max
## 40
## 41 max
## 42
## 43 max
## 44

```

Here we identify that the segment in the code that takes most time to run is the replicate function. This could be handled by some other primitive, or maybe the pre-allocation can be circumvented by having dynamic size of the vector.

Greedy knapsack solution

```

lineprof(greedy_knapsack(x = knapsack_objects[1:20000,], W = 3500))

```

```

## Reducing depth to 2 (from 47)

##      time alloc release dups                                ref
## 1 0.006 6.079          0 2256 c("compiler:::tryCmpfun", "tryCatch")
## 2 0.006 6.256          0 1000      c("stopifnot", "is.data.frame")
## 3 0.003 4.898          0   44      c("replicate", "sapply")
## 4 0.001 0.266          0    5      "order"
## 5 0.004 0.820          0   55      character(0)
##                                     src
## 1 compiler:::tryCmpfun/tryCatch
## 2 stopifnot/is.data.frame
## 3 replicate/sapply

```

```
## 4 order  
## 5
```

Not alot to improve on here.

Parallelizing brute force knapsack

The performance that could be gained is non-existent since the lapply used doesn't contain any calculations. So there is little to no sequential computations that are done. If we had a computationally heavy lapply segment then we could gain an decrease in computation time.