

Lab Report Knapsack

Simon Jönsson, Fanny Karelius

2017-10-11

Introduction

In this lab we implement different solutions for the **0/1-knapsack**- and **unbounded knapsack** problem. We have implemented both a parallel and non-parallel brute force solution, a dynamic programming solution and a solution using the greedy heuristic. We have documented the runtimes and the profiling of each solution.

Runtime of codes

Bruteforce knapsack solution

Parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:20,], W = 3500, parallel=FALSE))
```

```
##      user  system elapsed  
##    6.727    0.429    7.182
```

Non-parallel

```
system.time(brute_force_knapsack(x = knapsack_objects[1:20,], W = 3500, parallel=FALSE))
```

```
##      user  system elapsed  
##    5.540    0.381    5.938
```

Dynamic knapsack solution

```
system.time(knapsack_dynamic(x = knapsack_objects[1:500,], W = 3500))
```

```
##      user  system elapsed  
##    5.342    0.051    5.406
```

Greedy knapsack solution

```
system.time(greedy_knapsack(x = knapsack_objects[1:1000000,], W = 3500))
```

```
##      user  system elapsed  
##    2.383    0.158    2.550
```

Profiling

Bruteforce knapsack solution

Parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=TRUE))
```

```
## Reducing depth to 2 (from 12)
##      time alloc release dups                                ref
## 1 0.003 0.936      0 521                                c("::", "getExportedValue")
## 2 0.001 0.227      0 191                                c("parallel::mclapply", "::")
## 3 0.004 0.053      0 55                                c("parallel::mclapply", "lapply")
## 4 0.001 0.027      0 11 c("parallel::mclapply", "lazyLoadDBfetch")
## 5 0.001 0.005      0 0                                "parallel::mclapply"
## 6 0.002 0.046      0 9  c("parallel::mclapply", "selectChildren")
## 7 0.005 0.253      0 7  c("parallel::mclapply", "cleanup")
## 8 0.004 2.273      0 17                                c("lapply", "Filter")
##
##                                     src
## 1 ::/getExportedValue
## 2 parallel::mclapply/::
## 3 parallel::mclapply/lapply
## 4 parallel::mclapply/lazyLoadDBfetch
## 5 parallel::mclapply
## 6 parallel::mclapply/selectChildren
## 7 parallel::mclapply/cleanup
## 8 lapply/Filter
```

All segments of the code are in similar timesteps - quite tricky to identify bottlenecks. However one could look over using some primitive functions instead of using lapply to find the row with near-optimal value. A suggestions might be using **max()**.

Non-parallel

```
lineprof(brute_force_knapsack(x = knapsack_objects[1:12,], W = 3500, parallel=FALSE))
```

```
##      time alloc release dups                                ref      src
## 1 0.019 11.59      0 11691 c("lapply", "FUN") lapply/FUN
```

The lapply function might be exchanged with using a **max()** primitive. Generating the matrix with given parameters: weight and val one could use **max** to find the maximum val given that the weight $\leq W$.

Dynamic knapsack solution

```
lineprof(knapsack_dynamic(x = knapsack_objects[1:100,], W = 3500))
```

```
## Reducing depth to 2 (from 70)
##      time  alloc release dups                                ref
## 1 0.007  9.748  0.000 3626 c("compiler::tryCmpfun", "tryCatch")
## 2 0.100 153.462 143.548 1498                                c("matrix", "replicate")
## 3 0.001  2.399  0.000   0                                "matrix"
## 4 0.003 11.952  0.000 16386                                character(0)
## 5 0.001  1.859  0.000 13264                                "max"
## 6 0.001  4.092  0.000 3841                                character(0)
## 7 0.002  1.360  0.000 9985                                "max"
## 8 0.004  2.475  0.000 4765                                character(0)
## 9 0.001  0.631  0.000 1627                                "max"
## 10 0.003  1.901  0.000 4448                                character(0)
## 11 0.001  0.748  0.000 784                                "max"
## 12 0.005  2.660  0.000 5683                                character(0)
```

## 13	0.001	0.618	0.000	1358	"max"
## 14	0.002	0.786	0.000	2744	character(0)
## 15	0.001	0.399	0.000	158	"max"
## 16	0.004	1.882	0.000	3808	character(0)
## 17	0.001	0.618	0.000	905	"max"
## 18	0.004	2.282	0.000	5035	character(0)
## 19	0.002	0.876	0.000	1969	"max"
## 20	0.002	0.876	0.000	1807	character(0)
## 21	0.001	0.761	0.000	796	"max"
## 22	0.002	1.315	0.000	2797	character(0)
## 23	0.001	0.438	0.000	1494	"max"
## 24	0.002	1.224	0.000	1779	character(0)
## 25	0.002	1.521	0.000	3225	"max"
## 26	0.001	0.657	0.000	1573	character(0)
## 27	0.002	1.405	0.000	2878	"max"
## 28	0.001	0.748	0.000	1384	character(0)
## 29	0.002	1.141	0.000	2678	"max"
## 30	0.001	0.516	0.000	1226	character(0)
## 31	0.001	0.561	0.000	1065	"max"
## 32	0.002	0.593	0.000	1687	character(0)
## 33	0.001	0.894	0.000	696	"max"
## 34	0.006	3.320	0.000	7218	character(0)
## 35	0.001	0.683	0.000	1494	"max"
## 36	0.002	1.166	0.000	2504	character(0)
## 37	0.001	0.709	0.000	1318	"max"
## 38	0.003	1.630	0.000	3915	character(0)
## 39	0.002	0.786	0.000	1410	"max"
## 40	0.001	0.619	0.000	1132	character(0)
## 41	0.001	0.709	0.000	1278	"max"
## 42	0.011	8.508	80.594	20678	character(0)
##				src	
## 1	compiler::tryCmpfun/tryCatch				
## 2	matrix/replicate				
## 3	matrix				
## 4					
## 5	max				
## 6					
## 7	max				
## 8					
## 9	max				
## 10					
## 11	max				
## 12					
## 13	max				
## 14					
## 15	max				
## 16					
## 17	max				
## 18					
## 19	max				
## 20					
## 21	max				
## 22					
## 23	max				

```
## 24
## 25 max
## 26
## 27 max
## 28
## 29 max
## 30
## 31 max
## 32
## 33 max
## 34
## 35 max
## 36
## 37 max
## 38
## 39 max
## 40
## 41 max
## 42
```

Here we identify that the segment in the code that takes most time to run is the replicate function. This could be handled by some other primitive, or maybe the pre-allocation can be circumvented by having dynamic size of the vector.

Greedy knapsack solution

```
lineprof(greedy_knapsack(x = knapsack_objects[1:20000,], W = 3500))
```

```
## Reducing depth to 2 (from 18)

##      time alloc release dups                                ref
## 1 0.002 5.828          0 3215 c("compiler::tryCmpfun", "tryCatch")
## 2 0.006 1.219          0  41  c("stopifnot", "is.data.frame")
## 3 0.003 9.768          0  44  c("replicate", "sapply")
## 4 0.002 0.079          0   5  "order"
## 5 0.005 0.977          0  53  character(0)
##                                     src
## 1 compiler::tryCmpfun/tryCatch
## 2 stopifnot/is.data.frame
## 3 replicate/sapply
## 4 order
## 5
```

Not alot to improve on here.

Parallelizing brute force knapsack

The performance that could be gained is non-existent since the lapply used doesn't contain any calculations. So there is little to no sequential computations that are done. If we had a computationally heavy lapply segment then we could gain an decrease in computation time.